

Last Time:

- Hybrid methods for legged locomotion

Today:

- Iterative Learning Control
-

* What happens when our model has errors?

- Models are approximate
- Simpler models are often preferred even if they're less accurate
- Feed back (e.g. LQR/MPC) can often compensate for model errors
- Sometimes that isn't enough (e.g. very tight constraints, performance, safety).

* Several Options:

- 1) Parameter Estimation: Classical "System ID" / "gray-box" modeling. Fit e.g. masses, lengths, etc. in your model from data.

- + Very simple efficient
- + Generalizes well
- + Interpretable
- Assumes model structure

2) Learn Model: Fit generic function approximator to the full dynamics or residual. Classical "black-box" modeling / System ID

- + Doesn't assume model structure
- + Generalizes
- Not sample efficient. Requires lots of data

Improve the model

3) Learn a Policy: Standard model-free RL approach:
Optimize a function approximation of the controller

- + Makes few assumptions
- Doesn't generalize
- Not sample efficient, Requires lots of "rollouts"

4) "Transfer": Assume we have a reference trajectory computed with a nominal model. Improve it with data from the real system.

- + Makes few assumptions
- Assumes a decent prior model
- Doesn't generalize (task specific)
- + Very sample efficient

Improve the controller

* Iterative Learning Control (ILC)

- Can think of this as a very specialized policy-gradient method on the policy class:

$$u_n(\bar{u}_n) = \bar{u}_n - K_n(x_n - \bar{x}_n)$$

$\underbrace{\quad}_{\text{reference inputs}}$
 $\underbrace{\quad}_{\text{any tracking controller}}$

+ we only update \bar{u}_n

- Can think of this as SQP where we get the RHS vector from a rollout on the real system.
- Assume we have a reference trajectory \bar{x}, \bar{u} that we want to track:

$$\begin{array}{ll} \min_{\substack{x_1: N \\ u_1: N}} & J = \sum_{n=1}^{N-1} \frac{1}{2} (x_n - \bar{x}_n)^T Q (x_n - \bar{x}_n) + \frac{1}{2} (u_n - \bar{u}_n)^T R (u_n - \bar{u}_n) \\ & + \frac{1}{2} (x_N - \bar{x}_N)^T Q (x_N - \bar{x}_N) \end{array}$$

$$\text{s.t. } \dot{x}_{n+1} = f(x_n, u_n)$$

(other constraints)

- The KKT system for this problem looks like:

$$\begin{bmatrix} H & C^T \\ C & 0 \end{bmatrix} \begin{bmatrix} \partial z \\ \lambda \end{bmatrix} = \begin{bmatrix} -\nabla J \\ -C(z) \end{bmatrix}$$

where $\partial z = \begin{bmatrix} \partial x_1 \\ \partial u_1 \\ \vdots \\ \partial x_N \end{bmatrix}$, $C(z) = \begin{bmatrix} \vdots \\ f(x_n, u_n) - \bar{x}_{n+1} \\ \vdots \end{bmatrix}$, $C = \frac{\partial C}{\partial z}$

\uparrow
:
:
:

nominal dynamics model

$$H = \begin{bmatrix} Q & & \\ & R & \\ & & - \\ & & \\ & & Q_u \end{bmatrix} \quad \leftarrow \text{Gauss-Newton Hessian}$$

- Two important observations:
 - 1) If we do a rollout on the real system, $\mathcal{C}\bar{z} = 0$ always (for the true dynamics)
 - 2) Since we know J , given x_n, u_n from a rollout, we can compute ∂J
- Now we have RHS vector from the KKT system
- We also know H from the cost
- We can compute $C = \frac{\partial \mathcal{C}}{\partial z}$ using x_n, u_n and the nominal model
- However, since our nominal model is approximate and assuming x_n, u_n is already close to \bar{x}, \bar{u} , we can just use $C = \frac{\partial \mathcal{C}}{\partial z} |_{\bar{x}, \bar{u}}$, which can be computed offline
- Now just solve KKT system for Δz , update:
 $\bar{u} \leftarrow \bar{u} + \Delta u$, and repeat
- Can easily add inequality constraints (e.g. torque limits) and solve a QP.

* ILC Algorithms:

- Given nominal \bar{x}, \bar{u}

do :

$$x_{1:N}, u_{1:N} \leftarrow \text{rollout}(\bar{x}_0, \bar{u})$$

on real system

$$\delta x, \delta u \leftarrow \underset{\substack{\text{This is} \\ \text{a QP}}}{\text{argmin}} J(\delta x, \delta u)$$

$$\begin{aligned} \text{s.t. } \delta x_{n+1} &= A_n \delta x_n + B_n \delta u_n \\ u_{\min} &\leq u_n \leq u_{\max} \end{aligned}$$

$$\bar{u} \leftarrow \bar{u} + \delta u$$

while $\|x_n - \bar{x}_n\| \geq tol$

whatever you care about

* Why Should ILC Work?

- We've already seen approximations of Newton's method (e.g. Gauss-Newton)
- In general, these are called "quasi-Newton" or "inexact Newton" methods. Many variants e.g. BFGS / Newton-CG. Well developed theory.
- For a generic root-finding problem :

$$f(x + \delta x) \approx f(x) + \underbrace{\frac{\partial f}{\partial x} \delta x}_{\text{exact Newton step}} = 0$$

exact Newton step

- As long as $\mathcal{O}x$ satisfies:

$\|f(x) + J\mathcal{O}x\| \leq \eta \|f(x)\|$ for some $\eta < 1$,
an inexact Newton method will converge

- Convergence is slower than exact Newton

- This means we can use $J \approx \frac{\partial f}{\partial x} \neq \frac{\partial f}{\partial x}$ to compute $\mathcal{O}x$