

Last Time:

- How to drive a car

Today:

- Data-driven / model-free "behavioral" control
- Behavior Cloning

What if we don't have a model?

- Given input-output data, standard approach is to fit a model (System ID)
- Let's assume we have data u_{ref}, y_{ref} and we fit a linear state-space model:

$$x_u = Ax_u + Bu$$

$$y_u = Cx_u$$

- For an arbitrary change of coordinates $\tilde{x} = Q\tilde{x}$:

$$\tilde{x}_u = \underbrace{Q^{-1}AQ}_{\tilde{A}} \tilde{x}_u + \underbrace{Q^{-1}Bu}_{\tilde{B}}$$

$$y_u = \underbrace{CQ}_{\tilde{C}} \tilde{x}_u$$

\Rightarrow State definition is arbitrary!

* Can we just skip the state entirely?

A Different Take on Linear Systems:

- What does it mean to be linear?

$$f(ax_1 + bx_2) = af(x_1) + bf(x_2)$$

↑
vector
scalars

- "Superposition" \Rightarrow linear combination of inputs
 ↓
 " " " outputs

- For a linear dynamical system, superposition holds for trajectories!

$$\begin{bmatrix} y_1 \\ u_1 \\ y_2 \\ u_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} C \\ CI \\ CA \\ CA^2 \\ \vdots \\ CA^{n-1} \end{bmatrix} \underbrace{\begin{bmatrix} 0 & & & \\ & I & & \\ & & \ddots & \\ & & & I \end{bmatrix}}_{F} \begin{bmatrix} x_1 \\ u_1 \\ u_2 \\ \vdots \\ x_n \end{bmatrix}$$

$$\Rightarrow F(av^1 + bv^2) = afv^1 + bfv^2 = az^1 + bz^2$$

- Given enough trajectory data, we can span the space of all possible input-output trajectories of a given length:

$$z = az^1 + bz^2 + cz^3 + \dots$$

\Downarrow

$$= \begin{bmatrix} \vdots & \vdots & \vdots & \vdots \\ z^1 & z^2 & z^3 & \dots \\ \vdots & \vdots & \vdots & \vdots \end{bmatrix} \underbrace{\begin{bmatrix} a \\ b \\ c \\ \vdots \end{bmatrix}}_{w} \sim$$

H

- We can now use H in place of a state-space model (A, B, C) to do prediction + control.

* Example

- State prediction for damped SHO:

$$\ddot{q} = -Kq - C\dot{q} \Rightarrow \begin{bmatrix} \dot{q} \\ \ddot{q} \end{bmatrix} = \underbrace{\begin{bmatrix} 0 & 1 \\ -K & -C \end{bmatrix}}_A \underbrace{\begin{bmatrix} q \\ \dot{q} \end{bmatrix}}_x$$

$$Z = \begin{bmatrix} x_1 \\ \vdots \\ x_{50} \end{bmatrix} = H w$$

$$\min_w \frac{1}{2} \| E H w - x_1 \|_2^2 , \quad E = [I \quad 0 \quad 0 \quad \dots]$$

least-squares projection

- Closed-form solution:

$$w = (H^T E^T E H)^{-1} H^T E^T x_1 , \quad Z = H w$$

- We can do the same thing with only $y = Cx$

- We can also generate new dynamically feasible trajectories by sampling a random vector and projecting onto the data manifold ("de-noising")

Behavior Cloning:

- Let's assume we have input-output data from a closed-loop system being controlled by an "expert policy"

$$z' = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \end{bmatrix} \Rightarrow H = \begin{bmatrix} \cdot & \cdot & \cdot \\ z^1 & z^2 & \cdots \\ \vdots & \vdots & \vdots \end{bmatrix}$$

- Given some history of observations $y_{1:m}$, we can predict the rest of the input-output trajectory.
- We can get a feedback policy by taking clues from the predicted z :

$$\min_w \frac{1}{2} \| YHw - y_{1:m} \|^2, \quad Y = \begin{bmatrix} I & 0 & \cdots \\ 0 & I & 0 \cdots \end{bmatrix}$$

$$u_{mt+1} = VHw, \quad V = [0 \ 0 \ -I \ 0 \ \cdots \ 0]$$

Picks out $y_{1:m}$

* Example

- Cloning an LQR policy from random demonstrations

Data-Driven MPC :

- We can use the data matrix as our "dynamics model" in an MPC controller.
- Note that we need data with "sufficient excitation"? It can't be just closed-loop "on-policy" rollouts like in the prediction cloning case.

* Safe bet: inject noise into the systems during data collection

- Given a data matrix :

$$Z = \begin{bmatrix} y_1 \\ u_1 \\ y_2 \\ u_2 \\ \vdots \end{bmatrix} = Hw$$

- We can solve the following QP:

$$\min_{z, w} \frac{1}{2} z^T Q z + \frac{1}{2} \alpha w^T w \quad \text{regularizer on weights}$$

$$\text{s.t. } z = Hw$$

\leftarrow picks out first few outputs

$$y = Yz$$

\leftarrow Observations

$$\Rightarrow L = \frac{1}{2} z^T Q z + \frac{1}{2} \alpha w^T w + \lambda^T (Hw - z) + \mu^T (Yz - y)$$

$$\nabla L = 0 \Rightarrow \begin{bmatrix} Q & 0 & -I & Y^T \\ 0 & \alpha I & H^T & 0 \\ I & H & 0 & 0 \\ Y & 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} z \\ w \\ \lambda \\ \mu \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ y \end{bmatrix}$$

- Note that this is an output feedback controller (uses y instead of x). There's no separate state estimator.

What about nonlinear systems?

- In the linear case, the data manifold is a (hyper-) plane / subspace
- For nonlinear systems, it is a more complex curved manifold
 - ⇒ superposition doesn't hold anymore, but we can still define projection onto the manifold.
- = Diffusion methods essentially learn this projection